

Line Search Methods

Saurav Samantaray

Department of Mathematics

Indian Institute of Technology Madras

February 13, 2024



Line Search Method

- In each iteration of a line search method a search direction p_k is computed, and
- then its decided how far to move along that direction.

An iteration is given by

$$x_{k+1} = x_k + \alpha_k p_k \quad (1)$$

where $\alpha_k > 0$ (scalar) called the step length

. The success of a line search method depends on effective choices of both:

- 1 the direction p_k
- 2 the step length α_k

The Steepest Descent Direction

- The steepest descent direction $-\nabla f_k$ is the most obvious choice for search direction for a line search method.
- Among all Directions one could find from x_k , along $-\nabla f_k$, f decreases most rapidly.

Justification

- Consider any search direction p and step-length α , we have

$$f(x_k + \alpha p) = f_k + \alpha p^T \nabla f_k + \frac{\alpha^2 p^T \nabla^2 f(x_k + tp) p}{2} \text{ for some } t \in (0, \alpha)$$

- Let $\alpha \ll 1$ (small) and we consider the first-order approximation of f at $x_k + \alpha p$ around x_k as:

$$f(x_k + \alpha p) \approx f(x_k) + \alpha p^T \nabla f_k$$

- Change in f moving from x_k to $x_k + \alpha p$ is

$$f(x_k + \alpha p) - f(x_k)$$

The Steepest Descent Direction

- As the distance moved in the direction is α , therefore, the rate of change of f along the direction p at x_k is

$$\frac{f(x_k + \alpha p) - f(x_k)}{\alpha}$$

- which is coefficient of α , i.e.

$$p^T \nabla f_k$$

- This implies smaller the above value is, more descent can be achieved.
- Hence, the unit direction p of most rapid decrease is the solution to the problem

$$\min_p p^T \nabla f_k, \quad \text{subject to } \|p\| = 1.$$

The Steepest Descent Direction



$$p^T \nabla f_k = \|p\| \|\nabla f_k\| \cos \theta = \|\nabla f_k\| \cos \theta$$

where θ is the angle between p and ∇f_k .

- The minimiser is attained when $\cos \theta = -1$ and

$$p = -\frac{\nabla f_k}{\|\nabla f_k\|}$$




Figure illustrates this direction is orthogonal to the contours of the function

The Steepest Descent Direction

- At every step (iteration) in the steepest descent method the search direction is chosen along

$$p = -\nabla f_k$$

- α_k can be chosen in a variety of ways.
- One advantage of this method is it requires only the calculation of gradient (∇f_k), but not second derivatives.
- Line search methods may use search directions other than the steepest descent direction.

Descent Direction

Descent Direction

Any direction that makes an **angle of strictly less than $\frac{\pi}{2}$** radians with ∇f_k is guaranteed to produce a decrease in f , provided the step length is sufficiently small and is called a **descent direction**.

- Now consider

$$f(x_k + \epsilon p_k) = f(x_k) + \epsilon p_k^T \nabla f_k + \mathcal{O}(\epsilon^2).$$

- When p_k is a downhill (descent) direction, the angle θ_k between p_k and ∇f_k has $\cos \theta_k < 0$, so that

$$\begin{aligned} p_k^T \nabla f_k &= \|p_k\| \|\nabla f_k\| \cos \theta_k < 0 \\ \implies f(x_k + \epsilon p_k) &< f(x_k) \end{aligned}$$

- Most line search algorithms require p_k to be descent direction, because this property guarantees that the function

Newton Search Direction

- This direction is derived from the second-order Taylor series approximation to $f(x_k + p)$, which is

$$f(x_{k+p}) \approx f_k + p^T \nabla f_k + \frac{1}{2} p^T \nabla^2 f_k p \stackrel{\text{def}}{=} m_k(p). \quad (2)$$

- Assuming for the moment that $\nabla^2 f_k$ is positive definite, we obtain the Newton direction by finding the vector p that minimizes $m_k(p)$.
- By simply setting the derivative of $m_k(p)$ to zero, we obtain the following explicit formula:

$$p_k^N = -(\nabla^2 f_k)^{-1} \nabla f_k. \quad (3)$$

- The Newton direction is reliable when the difference between the true function $f(x_k + p)$ and its quadratic model $m_k(p)$ is not too large.

Newton Search Direction

- If $\nabla^2 f$ is sufficiently smooth, this difference introduces a perturbation of only $\mathcal{O}(\|p\|^3)$.
- Therefore, when $\|p\|$ is small

$$f(x_k + p) \approx m_k(p) \quad \text{quite accurately}$$

- The Newton direction can be used in a line search method when $\nabla^2 f_k$ is positive definite, as in this case we have

$$\begin{aligned} \nabla f_k^T p_k^N &= -p_k^{N^T} \nabla^2 f_k p_k^N \\ &< -\sigma_k \|p_k^N\|^2 \end{aligned}$$

for some $\sigma_k > 0$ (+ve definiteness of $\nabla^2 f_k$)

- Unless the gradient ∇f_k (and therefore the step p_k^N) is zero, we have

$$\nabla f^T p_k^N < 0$$

Newton Search Direction

- Unlike the steepest descent direction, there is a “natural” step length of 1 associated with the Newton direction.
- Adjust α only when it does not produce a satisfactory reduction in the value of f .
- Note that when $\nabla^2 f_k$ is not positive definite the Newton direction may not exist, since $(\nabla^2 f_k)^{-1}$ may not exist.
- Even when it is defined, it may not satisfy the descent property, and therefore is unsuitable.
- Methods that use Newton direction have fast rate of local convergence (**more on this later**).
- After a neighbourhood of the solution is reached, convergence to high accuracy often occurs in just a few iterations.

Newton Search Direction

- The main drawback is the need to calculate the Hessian $\nabla^2 f(x)$.
- Explicit computation of this matrix of second derivatives can sometimes be a cumbersome, error-prone, and expensive process.
- Finite-difference and automatic differentiation techniques come useful in avoiding the need to calculate second derivatives by hand.

Quasi-Newton Search Direction

- Quasi-Newton search directions provide an attractive alternative to Newton's method.
- They do not require computation of the Hessian and yet still attain a **superlinear** rate of convergence.
- In place of the true Hessian $\nabla^2 f_k$, they use an **approximation** B_k , which is updated after each step to take account of the additional knowledge gained during the step.
- The updates make use of the fact that changes in the gradient g provide information about the second derivative of f along the search direction.

Quasi-Newton Search Direction

- Using the Taylor's expansion we have

$$\nabla f(x + p) = \nabla f(x) + p^T \nabla^2 f(x) + \mathcal{O}(\|p\|^2)$$

- By setting $x = x_k$ and $p = x_{k+1} - x_k$, we obtain

$$\nabla f_{k+1} = \nabla f_k + \nabla^2 f_k (x_{k+1} - x_k) + \mathcal{O}(\|x_{k+1} - x_k\|^2).$$

Quasi-Newton Search Direction

- When x_k and x_{k+1} lie in a region near the solution x^* , within which $\nabla^2 f$ is positive definite, the final term in this expansion is eventually dominated by the $\nabla^2 f_k(x_{k+1} - x_k)$ term, and we can write.

$$\nabla^2 f_k(x_{k+1} - x_k) \approx \nabla f_{k+1} - \nabla f_k. \quad (4)$$

- We choose the new Hessian approximation B_{k+1} so that it mimics the property (4) of the true Hessian.
- That is we require it to satisfy the following condition, known as the secant equation:

$$B_{k+1}s_k = y_k, \quad (5)$$

where $s_k = x_{k+1} - x_k$, $y_k = \nabla f_{k+1} - \nabla f_k$.

Quasi-Newton Search Direction

Typically, additional conditions are imposed on B_{k+1} , such as

- Symmetry (Motivated by symmetry of the exact Hessian).
- a requirement that the difference between successive approximations B_k and B_{k+1} have low rank.

Two of the most popular formulae for updating the Hessian approximation B_k are:

Symmetric-rank-one (SR1) formula:

$$B_{k+1} = B_k + \frac{(y_k - B_k s_k)(y_k - B_k s_k)^T}{(y_k - B_k s_k)^T s_k}. \quad (6)$$

Quasi-Newton Search Direction

BFGS Formula:

$$B_{k+1} = B_k - \frac{B_k s_k s_k^T B_k}{s_k^T B_k s_k} + \frac{y_k y_k^T}{y_k^T s_k}. \quad (7)$$

- The difference between the matrices B_k and B_{k+1} is
 - rank-one for SR1.
 - rank-two in case of BFGS.
- Both updates satisfy the secant equation and both maintain symmetry.

Quasi-Newton Search Direction

- It can be shown that the BFGS formula generates positive definite approximations whenever the initial approximation B_o is positive definite and

$$s_k^T y_k > 0$$

- The quasi-Newton search direction is obtained by using B_k instead of the exact Hessian $\nabla^2 f_k$ in the Newton direction

$$p_k = -B_k^{-1} \nabla f_k.$$

- Some practical implementations of quasi-Newton methods avoid the need to factorise B_k at each iteration by updating inverse of B_k , instead of B_k itself.

Quasi-Newton Search Direction

- The equivalent formula for (SR1), applied to the inverse approximation

$$H_k :^{defn} = B_k^{-1} \quad \text{is}$$
$$H_{k+1} = (I - \rho_k s_k y_k^T) H_k (I - \rho_k y_k s_k^T) + \rho_k s_k s_k^T$$

where $\rho_k = \frac{1}{y_k^T s_k}$

- Then ρ_k is given by

$$\rho_k = -H_k \nabla f_k$$

- This matrix-vector multiplication is simpler than the factorisation / back-substitution produce.

Non-linear Conjugate Gradient Methods

- The direction here are generated by non-linear conjugate gradient methods.
- They have the form

$$p_k = -\nabla f(x_k) + B_k p_{k-1}$$

where B_k is a scalar that ensures that p_k and p_{k-1} are conjugate (to be defined later).

- Conjugate gradient methods were originally designed to solve systems of linear equations:

$$Ax = b$$

where the coefficient matrix A is symmetric and positive definite.

Non-linear Conjugate Gradient Methods

- The problem of solving this linear system is equivalent to the problem of minimising the convex quadratic function defined by

$$\phi(x) = \frac{1}{2}x^T Ax - b^T x.$$

- In general, non-linear conjugate gradient directions are much more effective than the steepest descent direction and are almost as simple to compute.
- These methods do not attain fast convergence rates of Newton or quasi-Newton methods.
- But, they have the advantage of not requiring storage of matrices.