

Large-Scale Unconstrained Optimization

Saurav Samantaray

Department of Mathematics

Indian Institute of Technology Madras

March 10, 2024



Introduction

- Many applications give rise to unconstrained optimization problems with thousands or millions of variables.
- Problems of this size can be solved efficiently only if the storage and computational costs of the optimization algorithm can be kept at a tolerable level.
- A diverse collection of large-scale optimization methods has been developed to achieve this goal, each being particularly effective for certain problem types.
- Some of these methods are straightforward adaptations of the methods described until now.
- Other approaches are modifications of these basic methods that allow approximate steps to be calculated at lower cost in computation and storage.

Introduction

- The non-linear conjugate gradient methods can be applied to large problems without modification, owing to its minimal storage demands and its reliance on only first-order derivative information.
- The Newton method in both line search and trust-region algorithms require matrix factorizations of the Hessian matrices.
- High quality software implementations are available, which are based on factorizations that can be carried out using sparse elimination techniques.
- Newton methods are plagued with issues related to computational cost and memory requirements of these sparse factorization methods.
- If the Hessian matrix can be formed explicitly, with the above problems sorted they constitute an effective approach for solving such problems.

Introduction

- Often, however, the cost of factoring the Hessian is prohibitive, and it is preferable to compute approximations to the Newton step using iterative linear algebra techniques.
- Inexact Newton methods that use these techniques, in both line search and trust-region frameworks have attractive global convergence properties and may be super-linearly convergent for suitable choices of parameters.
- There are variants of the quasi-Newton approach, which use Hessian approximations that can be stored compactly by using just a few vectors of length n .
- These methods are fairly robust, inexpensive, and easy to implement, but they do not converge rapidly.

Inexact Newton Methods

- The basic Newton step p_k^N is obtained by solving the symmetric $n \times n$ linear system

$$\nabla^2 f_k p_k^N = -\nabla f_k. \quad (1)$$

- p_k^N can be obtained by solving the above equation (1) approximately, via inexpensive iterative solvers.
- For example the conjugate gradient (CG) method can be employed to get p_k^N .
- Both line search and trust region approaches can be derived based on this approximation, which falls in the general family of inexact Newton methods.
- In addition, we can implement these methods in a Hessian-free manner, so that the Hessian $\nabla^2 f_k$ need not be calculated or stored explicitly at all.

LOCAL CONVERGENCE OF INEXACT NEWTON METHODS

- Consider the residual for the equation (1) as:

$$r_k = \nabla^2 f_k p_k + \nabla f_k \quad (2)$$

where p_k is the inexact Newton step.

- The CG iterations are terminated when

$$\|r_k\| \leq \eta_k \|\nabla f_k\|, \quad (3)$$

where the sequence $\{\eta_k\}$ (with $0 < \eta_k < 1$ for all k) is called the forcing sequence.

LOCAL CONVERGENCE OF INEXACT NEWTON METHODS

Theorem

Suppose that $\nabla^2 f(x)$ exists and is continuous in a neighbourhood of a minimizer x^* , with $\nabla^2 f(x^*)$ is positive definite. Consider the iteration $x_{k+1} = x_k + p_k$ where p_k satisfies (3), and assume that $\eta_k \leq \eta$ for some constant $\eta \in [0, 1)$. Then, if the starting point x_0 is sufficiently near x^* , the sequence $\{x_k\}$ converges to x^* and satisfies

$$\|\nabla^2 f(x^*)(x_{k+1} - x^*)\| \leq \hat{\eta} \|\nabla^2 f(x^*)(x_k - x^*)\|, \quad (4)$$

for some constant $\hat{\eta}$ with $\eta < \hat{\eta} < 1$.

LINE SEARCH NEWTON–CG METHOD

- In the line search Newton–CG method, also known as the truncated Newton method, the search direction is computed by applying the CG method to the Newton equations:

$$\nabla^2 f_k p_k^N = -\nabla f_k;$$

- and attempt to satisfy a termination test of the form

$$\|r_k\| \leq \eta_k \|\nabla f_k\|,$$

- The CG method is designed to solve positive definite systems.
- However, the Hessian $\nabla^2 f_k$ may have negative eigenvalues when x_k is not close to a solution.
- The CG iteration is terminated as soon as a direction of negative curvature is generated.

LINE SEARCH NEWTON–CG METHOD

- This adaptation of the CG method produces a search direction p_k that is a descent direction.
- Moreover, the adaptation guarantees that the fast convergence rate of the pure Newton method is preserved, provided that the step length $\alpha_k = 1$ is used whenever it satisfies the acceptance criteria.
- For purposes of this algorithm rewrite the linear system (1) in the form

$$B_k p = -\nabla f_k \quad (5)$$

where B_k represents $\nabla^2 f_k$.

- For the inner CG iteration, denote the search direction by d_j and the sequence of iterates that it generates by z_j .

LINE SEARCH NEWTON–CG METHOD

- When B_k is positive definite, the inner iteration sequence $\{z_j\}$ will converge to the Newton step p_k^N that solves (5).
- At each major iteration, a tolerance ε_k that specifies the required accuracy of the computed solution.
- For concreteness the forcing sequence is chosen to be $\eta_k = \min(0.5, \sqrt{\|\nabla f_k\|})$ to obtain a super-linear convergence rate (one may choose differently as well).

Line Search Newton–CG

Algorithm

```

Given initial point  $x_0$ ;
for  $k = 0, 1, 2, \dots$ 
  Define tolerance  $\varepsilon = \min(0.5, \sqrt{||\nabla f_k||}) ||\nabla f_k||$ ;
  Set  $z_0 = 0$ ,  $r_0 = \nabla f_k$ ,  $d_0 = -r_0 = -\nabla f_k$ ;
  for  $j = 0, 1, 2, \dots$ 
    if  $d_j^T B_k d_j \leq 0$ 
      if  $j = 0$ 
        return  $p_k = -\nabla f_k$ ;
      else
        return  $p_k = z_j$ ;
    set  $\alpha_j = r_j^T r_j / d_j^T B_k d_j$ ;
    Set  $z_{j+1} = z_j + \alpha_j d_j$ ;
    Set  $r_{j+1} = r_j + \alpha_j B_k d_j$ ;
    if  $||r_{j+1}|| < \varepsilon$ 
      return  $p_k = z_{j+1}$ ;
    Set  $\beta_{j+1} = r_{j+1}^T r_{j+1} / r_j^T r_j$ ;
    Set  $d_{j+1} = -r_{j+1} + \beta_{j+1} d_j$ ;
  end (for)
  Set  $x_{k+1} = x_k + \alpha_k p_k$ , where  $\alpha_k$  satisfies the Wolfe, Goldstein, or
  Armijo backtracking conditions (using  $\alpha_k = 1$  if possible);
end
  
```

LINE SEARCH NEWTON-CG METHOD

- The main differences between the inner loop of the above algorithm and the original CG are that the specific starting point $z_0 = 0$ is used;
- and the use of a positive tolerance ε_k allows the CG iterations to terminate at an inexact solution;
- and the negative curvature test $d_j^T B_k d_j \leq 0$ ensures that p_k is a descent direction for f at x_k .
- If negative curvature is detected on the first inner iteration $j = 0$, the returned direction $p_k = -\nabla f_k$ is both a descent direction and a direction of non-positive curvature for f at x_k .